

The Risks of Using a General-Purpose Commercial LLM in Judicial Support Applications: The Case of ChatGPT in Argentine Courts

Mahmoud Khundaqjie 2854994

Bo Van Laethem 2839082

Project Group 7

1. Introduction

In the 1967 short story “I Have No Mouth and I Must Scream” by Harlan Ellison, the story revolves around an AI system called AM that holds abject hatred towards humanity. While there is no system like that in reality, at least not yet, it poses the important challenge of regulating the use of artificial intelligence to avoid such a situation.

Many problems still plague AI today, hallucination [1], for example, is but one of the many problems that AI in a legal context would bring about. But are there any actual implementations of AI in the legal world? Yes, there are, and among the most prominent is Argentina using AI in many forms, namely the commercially available ChatGPT is being used by prosecutors and judges on an individual, unregulated basis while also being the basis of more specialized systems, all of which are support tools to increase efficiency [2], but this efficiency comes at a big risk, which we aim to determine by answering the research question.

What are the risks of Argentina using a general-purpose commercial LLM in judicial support applications?

2. Domain

The Argentinian government has implemented and experimented on the use of numerous AI systems built on ChatGPT or other commercial AI chatbots such as Google’s Gemini for the express purpose of improving efficiency. This is evidenced by the fact that one of the main purposes of these systems is stated to be summarization for assisting judicial decisions. One such example is the experimental Relmo system, which is a small, ChatGPT-powered chatbot the Supreme Court tested, which is designed to analyse and summarise court rulings for judicial use [2].

One example of a use of ChatGPT on an individual level in Argentina was a judge using it to make draft a more simply worded and accessible version of a paragraph [2]. An example of an instated system based on ChatGPT is the Genaro system, which is given a PDF of a case ruling draft and outputs the strengths and weaknesses of said draft. Importantly, this system cannot filter out data from inadequate citations.

3. Background & Literature Review

3.1 AI in Legal Practice

AI in its current state is already quite deeply embedded, directly and indirectly, in the legal world. Indirect implementation is namely in tools already in use, such as Gemini in Chrome or Copilot in Windows laptops, to implementation in legal databases such as Westlaw [3].

Beyond that, many legal professionals use AI tools on an individual basis, just as the prosecutors and judges in Argentina do, usually for the purposes of legal research, document review and drafting documents [3].

The main problem with AI in the legal world is that these commercial tools, whether it be ChatGPT, Gemini, Copilot or anything else available to the public, is that these tools were not built for the purpose of legal use.

Of the commercial tools, some have explicitly prohibited doing so in their terms of service, such as ChatGPT with giving advice that would usually be given by legal professionals being prohibited in the OpenAI usage policies [4].

3.2 AI Use in South American Courts

3.2.1 Brazil's MARIA

The MARIA system is a purpose-built LLM, developed in collaboration with the Elogroup consulting Ltda., and was implemented by the Brazilian Supreme Court in December of 2024, with the purpose of relieving the Brazilian courts' backlog of nearly 84 million cases. It does three tasks: automatically generates summaries from judges' votes, drafts procedural reports for special appeals, and conducts preliminary reviews of incoming complaints [5].

The key contrast between this system and using the commercially available ChatGPT is that every single document generated by MARIA is reviewed by legal professionals, while the Court's Technology and Innovation Secretariat oversees the process, which is legitimate human oversight, unlike the individual basis in which ChatGPT is used in Argentina and other places across the world.

However, despite this human oversight and the fact that MARIA does not make any decisions, its contributions can subtly influence outcomes, which is indicative of the fact that even purpose-built systems can come with their own risks.

4. Computational Solution

4.1 How ChatGPT Works

4.1.1 Inputs and Outputs

The way ChatGPT processes text and other inputs is done in a very systematic way involving multiple steps which work together to generate a desirable output. These steps include tokenization, embedding, transformer-based processing, and output generation [6].

In the first step, tokenization, the user input is converted to a form that can be processed by the model. In order for the model to process it, input text is split into smaller parts known as tokens. Tokens can refer to any form, including words, word fragments, punctuation marks, or characters.

These tokens are then encoded to create embeddings. This step is done by transforming the tokens into numerical vectors containing information about the semantics and syntax of the text.

These embeddings are then sent to ChatGPT's transformer architecture for processing done through the use of a self-attention mechanism, which helps the model to assess the relationship between each token and all other tokens derived from the input, building meaning and contextual understanding.

ChatGPT then employs patterns that it has learned during its training period in order to predict the token that is the most probable to follow in the given context. This prediction procedure is iterated for each new token until a stop condition is reached, where the output is converted to human-readable text and presented to the user.

4.1.2 Data Storage and Retention


When ChatGPT is used, the user's interactions (their prompts, responses, uploaded files, images, etc...) are stored in OpenAI's servers. This information is used to maintain security, prevent abuse, troubleshoot issues and improve system performance [7]. The user can, however, disable the option of using their data for model training through the platform's data controls without hindering their user experience.

Data retention settings vary from user to user depending on what plan or mode you are using. When using the Temporary Chat feature of ChatGPT, user interactions are not stored in memory and are therefore not used for training purposes. As for those who use consumer versions of ChatGPT (like free or plus), user data usage in training is turned on by default and conversations are generally stored in the user's chat history until they are manually deleted.

When a conversation is deleted, it is removed from the user's account immediately and is typically scheduled for permanent deletion from OpenAI's systems within 30 days, unless legal, security, or regulatory requirements require longer retention [8].

On the other hand, for business services such as ChatGPT Enterprise, ChatGPT Edu, and the OpenAI API, user data is not used in training by default. Additionally, Enterprise users have additional controls over data retention, privacy, and storage policies depending on their service agreement [8].

4.2 System Pipeline Flowchart



Flow Chart: ChatGPT judicial usage in Argentina

First, a judicial employee drafts a ChatGPT prompt, likely summarizing facts or asking legal questions.

In the next step, the process checks whether the prompt includes personally identifiable information (PII). If this is the case, then the system flags it as a privacy concern. A possible risk here is the potential leakage of sensitive court data to a third-party system beyond the court's control.

After that, the court chooses whether the interaction happens through ChatGPT's consumer interface or the OpenAI API.

Afterwards, the prompt is sent over an encrypted HTTPS connection, and a problem here is that encryption protects transit, but not necessarily downstream storage, logging, policy compliance, or misuse of submitted data.

Next, OpenAI receives and routes the request. Here data is logged for abuse monitoring, and, even if intended for safety, logs can contain sensitive judicial material unless minimized or redacted.

Then the generated answer is sent back to the court and the conversation is stored, making judicial prompts and outputs become retained records outside the official court system.

Next, a human should review the response before use. If this isn't done, the system flags this as a TOS/policy problem and the system restarts. If reviewed, the human checks for hallucinated law, facts, citations, or reasoning and either rejects or accepts it.

If the output is rejected, and the prompt is revised. A possible risk, however, is that hallucinations can be subtle and may survive review if staff lack time or expertise.

If the output is accepted, it becomes part of the judicial decision.

Finally, the ruling enters the official court record. By this point, any privacy failure, unchecked hallucination, or overreliance on AI has become part of the judicial system rather than just a drafting issue.

5. Evaluation

5.1 Evaluation Dimensions

- Hallucination — do cited sources exist and support the claims?
- Legal Accuracy — does ChatGPT's output match real case outcomes?
- Consistency — does repeated input produce the same output?

5.2 Experimental Procedure, Variables and Legal Concerns of the Procedure

5.2.1 Experimental Procedure

A selection of anonymized Argentine court cases will be loaded into ChatGPT and evaluated across the three evaluation dimensions using identical prompts across four trials per case, with each output scored by legal experts.

To test hallucination, citations produced by ChatGPT will be verified under two criteria: does the source exist? And does the source support the claim made?

To test legal accuracy, ChatGPT's generated output for each case is compared against the real outcome by qualified legal experts who will score the system on framing, outcome and reasoning.

To test consistency, each case will be run 4 times under identical, then each of the 4 runs are compared to each other and scored accordingly.

To address any legal concerns with using case data, all the information will be anonymized. Names and any personal information based on the definition of the GDPR will be redacted.

5.2.2 Variables

The variables of the experiments are as follows: the independent variable is the cases input for 4 trials, as multiple cases will be used and this is the one variable in the experimenter's control. The dependent variables are the three evaluation dimensions: hallucination score, legal accuracy score and consistency score, as these are the variables that the experts will be analyzing. The controlled variables, which are kept the same in each trial of each case, are the version of ChatGPT used, the prompt used and the legal experts who will analyze the cases.

5.3 Scoring System

Dimension	Max Points	Scoring
Hallucination	40	Citation exists and supports claim: 0 deducted; Citation exists but doesn't support claim: -5; Citation doesn't exist: -10
Legal accuracy	35	Correct reasoning: 0/7.5/15; Correct outcome: 0/5/10; Appropriate framing: 0/5/10
Consistency	25	All 4 runs agree: 25; 3/4 agree: 18; 2/4 agree: 10; No consistency: 0
Total	100	85–100 reliable · 70–84 needs oversight · 50–69 high risk · <50 unfit

5.4 Internal vs External Evaluation

To achieve the best possible evaluative results, there needs to be two types of evaluations done on ChatGPT in the judicial realm, internal evaluation and external evaluation.

Internal evaluation refers to the structured experiment and scoring system outlined in 5.2 and 5.3 respectively, while external evaluation refers to the independent assessment of ChatGPT's outputs by qualified Argentine legal professionals who were not involved in designing the experiment.

These experts evaluate whether the outputs are legally sound, contextually appropriate, and consistent with Argentine legal standards, without reference to the internal scoring rubric.

6. Legal and Ethical Considerations

6.1 Legal Considerations

6.1.1 EU AI Act – Article 6 & Annex III(a): High-Risk Classification

While Argentina is not an EU nation, analyzing the implementation of ChatGPT based on the EU AI Act is still relevant, as it is the first AI related regulation in the world. The AI Act is a risk-based regulation that classifies AI into 4 categories, prohibited actions, high risk, minimal risk and low risk [9].

According to Article 6 and Annex III(a) of the AI Act, using an AI in a legal context would be considered a high-risk use of AI. Again, while Argentina is not an EU member state, this classification sets the basis for an argument that using AI in a legal context is highly risky and that the obligations that come with this classification are not met [9].

6.1.2 GDPR – Article 9(1): Special Category Data

Once again, Argentina is not an EU member state, and so the GDPR does not officially apply to them, but it serves as a good referential framework.

GDPR Article 9 outlines many types of data that are considered special case data that are illegal to process unless legitimate bases are established. Among those types is legal data, specifically legal data relating to criminal convictions and offenses, which by using ChatGPT is sent and stored in OpenAI servers [7].

Thus, it can be said that under the GDPR using case data, even individually in the case of the Argentinian judges and prosecutors, is illegal as they use it for ongoing cases that can be influenced and that may not be anonymized, and, though they do not have to abide by the GDPR, it still suggests that their actions may be harmful and should be scrutinized [10].

6.1.3 OpenAI Usage Policy: Prohibition on Legal Decision-Making

Since Argentina's court uses ChatGPT on an individual level, it binds it to OpenAI's usage policies, analyzing their use of ChatGPT through OpenAI's usage policies reveals risks that OpenAI itself has identified when its systems are used in legal contexts. As their usage policies clearly prohibit the use of their systems in a setting that requires a license, unless a licensed practitioner is involved, additionally, OpenAI emphasizes that their systems should not be used to make high-stakes decisions without meaningful human oversight [4].

This creates a contradiction when examining the documented use of ChatGPT within Argentine courts. Although judges and prosecutors often claim that AI-generated text is reviewed by humans before incorporation into rulings, the system is still being used for tasks such as legal reasoning, summarising evidence, drafting judicial language, and influencing decision-making processes. By looking at the flowchart, it is evident that ChatGPT outputs may ultimately become incorporated into judicial rulings. Therefore, making AI participate in a process that directly affects the legal rights of individuals, which is the type of legal activity that OpenAI identifies as requiring strict human oversight and caution.

Thus, evaluating Argentina's use of ChatGPT through OpenAI's own standards reveals a violation of OpenAI's usage policies.

6.2 Ethical Considerations

In the context of ethics and ethical considerations there is a lot to consider when it comes to using commercially available AI technologies in the legal world.

Firstly, there is the risk of having sensitive data present in case files be shared to the AI provider's databases, and, while it is not used, it is stored by the companies, as stated in OpenAI's memory FAQs. On top of the fact that this is not allowed by the GDPR, the invasion of the people involved in the cases' privacy is a significant ethical concern. [8]

Secondly, there is the harm that could be done to people whose ongoing cases are put through AI systems, and the damage that could be done to them should AI misconstrue or misrepresent their case, be it through hallucination or biased training data. At best, it could be a slightly increased sentence and at worst it could be wrongful incrimination. This is something that cannot and should not be ignored.

7. Conclusion

7.1 Answer to the Research Question

In conclusion, the answer to the question of “**What are the risks of Argentina using a general-purpose commercial LLM in judicial support applications?**” is that there are many risks, but also some benefits. Of these benefits are mainly computational power from using a server-hosted LLM and the improvement in efficiency that using it may have on the legal process, but the risks are very much a large concern. The first concern is the concern of automation bias. This is caused by issues such as hallucination or bias which may cause judges to deliver false judgements. The other major concern is privacy where the information in case files being uploaded to an external server is concerned.

7.2 Limitations & Future Work

The main limitations of this paper and the research we conducted were the lack of access to legal materials and lack of contact with real legal experts, which made it so we could only propose a hypothetical experiment to test ChatGPT in the judicial system of Argentina. Future work should not only perform the experiment, or design their own, but it could also be more focused, such as answering other questions such as, how can we regulate commercial AI in the legal space? Or are purpose-built LLMs better and safer?

References

- [1] Nature. (2025). Risk of hallucination in LLMs.
<https://www.nature.com/articles/s41599-025-05924-3>
- [2] Tech and Justice, BSG Oxford. (n.d.). Argentina.
<https://www.techandjustice.bsg.ox.ac.uk/research/argentina>
- [3] Canadian Bar Association. (n.d.). Use of AI in legal practice: An overview.
<https://cba.org/resources/practice-tools/ethics-of-artificial-intelligence-for-the-legal-practitioner/2-use-of-ai-in-legal-practice-an-overview/>
- [4] OpenAI. (n.d.). Usage Policies. <https://openai.com/policies/usage-policies/>
- [5] SSRN. (2025). MARIA – Brazil’s judicial AI system.
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5146151
- [6] OpenAI. (n.d.). Structured Outputs – API Documentation.
<https://developers.openai.com/api/docs/guides/structured-outputs>
- [7] OpenAI. (n.d.). Data Usage for Consumer Services FAQ.
<https://help.openai.com/en/articles/7039943-data-usage-for-consumer-services-faq>
- [8] OpenAI. (n.d.). Memory FAQ. <https://help.openai.com/en/articles/8590148-memory-faq>
- [9] European Parliament & Council of the European Union. (2024). Regulation (EU) 2024/1689 – Artificial Intelligence Act. Official Journal of the European Union.
<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>
- [10] European Parliament & Council of the European Union. (2016). Regulation (EU) 2016/679 – General Data Protection Regulation. Official Journal of the European Union.
<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32016R0679>